

TABLE S2: TWFE ESTIMATES

	(1)
$Treat_i \times Post_t$	0.0021*** (0.0003)
Const	0.0084*** (0.0000)
Unit Fixed Effect	Y
Year Fixed Effect	Y
Observations	4,751,559
No. clusters	446,725
R-squared	0.3877

Table S2. TWFE estimates of the exposure effect via review invitation on downstream citations to corresponding authors. This table reports estimates from our baseline TWFE difference-in-differences specification, restricted to reviewer invitees in Click-Decline and Auto-Decline groups. The dependent variable is the reviewer invitees' average annual number of citations to the focal corresponding author. Model (1) shows that reviewer invitees who *Click-Declined* an invitation cite the corresponding author 0.0021 additional times per year in the post-invitation period relative to reviewer invitees who *Auto-Declined* ($p = 0.000$). This estimate represents roughly a 14% increase over the post-invitation citation rate of the Auto-Decline group. The regression includes reviewer–manuscript fixed effects and calendar-year fixed effects, and excludes singleton observations, thereby controlling for all time-invariant pair characteristics and common year-specific shocks. Robust standard errors clustered at the reviewer invitee level are reported in parentheses. (***) $p < 0.01$, (**) $p < 0.05$, (*) $p < 0.1$).

To ensure the robustness of our event-study results, we update the TWFE specification by replacing the single $Treat_i \times Post_t$ term with interactions between treatment status and a set of event-time dummies. The results, shown in Figure S1, confirm the parallel trends assumption and are consistent with our main findings.




Figure S1. Event-study estimates from dynamic TWFE. Estimates include relative-year indicators interacted with treatment status, un-interacted event-time dummies, and the same fixed effects used in the static TWFE model. Each marker plots β_k from the event-study specification with fixed effects, measuring the difference in the average yearly number of citations a reviewer invitee gives the focal corresponding author (CA) k years relative to the invitation with one year before omitted as the baseline. Bars show the 95% confidence interval derived from the regression standard errors. Pre-treatment coefficients hover around zero. Post-invitation effects turn positive and grow steadily, reaching roughly 0.003 additional citations per year five years later.

S3.2 Staggered Difference-in-Differences

In our setting, reviewer invitees are exposed to treatment at different points in time, and the duration of post-treatment observation varies across individuals. For estimates from both the canonical Difference-in-Differences and traditional two-way fixed effects models to be interpreted as valid average treatment effects, they require the assumption of homogeneous treatment effects across cohorts and over time. However, as demonstrated by Goodman-Bacon [2], TWFE estimators aggregate all possible two-by-two comparisons—including comparisons between early- and late-treated groups—which can introduce negative weights and yield biased estimates in the presence of treatment effect heterogeneity.

To address this concern, we implement the CSDiD estimator proposed by Callaway and Sant’Anna [3]. This method improves *ATT* (i.e., average treatment effect on the treated) estimation under staggered treatment timing by estimating group-time average treatment effects and appropriately weighting them. Unlike TWFE, CSDiD avoids the problematic implicit

weighting structure and produces unbiased estimates even when treatment effects vary across cohorts or time. It estimates a cohort- and time-specific average treatment effect:

$$ATT(g, t) = E[Y_{it}(1) - \widetilde{Y}_{it}(0) \mid G_i = g]$$

where g denotes the year when a unit is first treated, t is the calendar year, $Y_{it}(1)$ is the observed outcome under treatment, and $\widetilde{Y}_{it}(0)$ is the counterfactual outcome had the unit remained untreated. This effect is defined only for post-treatment periods ($t \geq g$). We estimate $ATT(g, t)$ with *CSDiD*, which delivers three ways to collapse the cell-level effects which are reported in Table S3:

- Simple: Each cohort-time cell's $ATT(g, t)$ is weighted by the number of observations in that cell, and the overall effect is the sample-size-weighted mean across all g and t .
- Cohort-weighted: Within every treated cohort g , first take a sample-size-weighted average of its $ATT(g, t)$ across post-treatment periods; then weight those cohort-level means by the total number of observations in each cohort.
- Calendar-weighted: For each post-treatment calendar year t , average $ATT(g, t)$ across cohorts using sample-size weights; then combine the year-specific means by weighting them according to the number of treated observations appearing in that year.

By estimating effects cohort-by-cohort and then averaging, the *CSDiD* estimator sidesteps the contamination bias that plagues conventional two-way fixed-effects models under staggered adoption and permits flexible reporting of dynamic or cohort-specific treatment patterns.

Model (1) of Table S3 applies the simple Callaway–Sant’Anna aggregation and yields an ATT of 0.0011 ($p = 0.000$), which reflects that Click-Denial reviewer invitees cite corresponding authors 0.0011 more times per year than Auto-Denial reviewer invitees after the invitation. While modest in absolute terms, this still reflects a 7% relative increase compared to the Auto-Denial group’s post-invitation citation rate. The cohort-weighted (Model 2) and calendar-weighted (Model 3) specifications similarly confirm a significant and consistent positive effect. Additionally, the event study estimates in Figure S2 show that the increase begins in the year of the invitation and persists over time.

TABLE S3: STAGGERED DID ESTIMATES

	(1)	(2)	(3)
ATT	0.0011*** (0.0004)	0.0007** (0.0003)	0.0036* (0.0019)
Observations	4,656,924	4,656,924	4,656,924
Type	simple	cohort	calendar

Table S3. Average treatment effects from staggered DiD specifications. This table reports the average treatment effect on the treated (ATT) estimated with the Callaway-Sant’Anna staggered difference-in-differences approach under three aggregation schemes: simple (Model 1), cohort-weighted (Model 2), and calendar-weighted (Model 3). Across all three aggregation schemes—simple (0.0011, $p = 0.002$), cohort-weighted (0.0007, $p = 0.015$), and calendar-weighted (0.0036, $p = 0.061$)—the ATT remains positive and statistically significant, indicating a consistently robust treatment effect. Robust standard errors clustered at the reviewer invitee level are shown in parentheses. (** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$).




Figure S2. Staggered DiD event-study coefficients. The plot shows event-specific *ATT* estimates from the Callaway–Sant’Anna estimator, first averaged within each relative year t and then aggregated across cohorts with cohort-size weights. Each marker represents the estimated change in a reviewer invitee’s yearly citations to the focal corresponding author (CA) k years relative to the invitation. Whiskers depict 95% confidence intervals based on the *CSDiD* robust standard errors. Coefficients for pre-treatment years fluctuate around zero, supporting parallel trends, whereas post-invitation effects become increasingly positive, reaching roughly 0.006 additional citations per year five years after treatment. Note that $k = -10$ is not reported because *CSDiD* forms each pre-treatment contrast relative to the immediately prior year, and the window stops at -10 .

S4 Supplementary Results for Instrumental Variable Estimation

We adopt a two-stage least-squares (2SLS) approach for instrumental variable estimation. The instrument is a binary indicator for whether the invitation was sent on a weekend, and the unit of

analysis remains the reviewer–manuscript pair. Beginning with the same panel used in our DiD estimations, we collapse each reviewer–manuscript trajectory to two observations—the mean outcome over the entire pre-invitation window and the mean outcome over the post-invitation window—before running the IV procedure.

In the first-stage regressions, the weekend-assignment instrument is exceptionally strong: the Kleibergen–Paap rk F-statistic is 1,275.76 in the model without control variables and 1,220.48 in the specification that includes control variables—both far exceeding the Stock & Yogo’s [4] 10% weak-instrument critical value of 16.38, safeguarding the 2SLS estimates against weak-instrument bias.

TABLE S4: FIRST-STAGE RESULTS OF IV TEST

	(1)	(2)
<i>IV_weekend</i>	0.0469*** (0.0013)	0.0457*** (0.0013)
# Collaboration for Both		-0.0004* (0.0002)
# Citation Reviewer to Author		0.0003*** (0.0001)
# Citation Author to Reviewer		-0.0002*** (0.0001)
Reviewer # Publications (log)		0.0077*** (0.0004)
Reviewer H-index (log)		-0.0022*** (0.0004)
Reviewer # Pub on Journal		-0.0032*** (0.0002)
Reviewer # Ref to Journal		0.0008*** (0.0000)
Const	0.1837*** (0.0012)	0.1463*** (0.0021)
Kleibergen–Paap rk F-statistics	1275.76	1220.48
Observations	742,229	742,229
No. clusters	450,773	450,773
R-squared	0.0010	0.0053

Table S4. First-stage regression from the 2SLS instrumented analysis. This table reports first-stage estimates from a two-stage least squares (2SLS) regression model in which the binary instrument *IV_weekend* indicates whether the invitation was sent on a weekend. Model (2) controls for prior collaboration with the corresponding author, reviewer invitees’ research experience, and ties to the journal. The Kleibergen–Paap rk