

Table 1: **Evaluation of LLM extraction accuracy on author–affiliation–email triplet consistency.** We report the structural mismatch rates for affiliations and emails, the unparseable output rate (δ_{parse}), and the overall success rate across five GLM family models evaluated on the metadata benchmark.

Model	δ_{aff} (%)	δ_{email} (%)	δ_{parse} (%)	Success Rate (%)
glm-4-plus	5.01%	4.94%	0.81%	86.82%
glm-4-air	49.98%	17.11%	0.51%	44.73%
glm-4-flash	76.39%	43.27%	0.62%	18.52%
glm-3-turbo	76.07%	32.34%	1.34%	20.90%

Table 2: **Cross-venue token usage and affiliation disclosure (2021–2025).** Token consumption statistics and available structured affiliation metadata across major AI/ML venues.

Venue Metadata			Affiliation Disclosure			Token Consumption		
Venue	Years	#Papers	Author–Affil.	Affil.–Country	% Multi-Affil.	Prompt Tokens	Completion Tokens	Total Tokens
AAAI	2021–2025	11347	Partial	No	38.10%	1.80M	19.22M	21.02M
ACL	2021–2024	5724	No	No	39.47%	0.88M	8.86M	9.74M
CVPR	2021–2024	8803	No	No	38.12%	1.36M	12.68M	14.03M
EMNLP	2021–2024	7365	No	No	32.94%	1.12M	11.35M	12.47M
ICCV	2021,2023	3768	No	No	34.95%	0.57M	5.44M	6.01M
ICML	2023–2024	4438	No	No	39.12%	0.63M	7.00M	7.62M
IJCAI	2021–2024	3484	No	No	34.67%	0.53M	5.75M	6.28M
NeurIPS	2021–2024	14190	No	No	31.28%	2.09M	17.35M	19.44M

Note: “Author–Affiliation” refers to structured author–affiliation pair disclosure in submission metadata. “Partial” indicates that only one affiliation per author is disclosed, even when multiple affiliations are listed in the camera-ready version. No venues disclose structured affiliation–country information. Multi-affiliation percentages are estimated from parsed camera-ready PDFs using our automated extraction pipeline enabled by *glm-4-plus*.

we define the affiliation mismatch $\delta_{\text{aff}}^i = \mathbf{1}(\mathcal{A}^i, a^i)$ and email mismatch $\delta_{\text{email}}^i = \mathbf{1}(e^i, a^i)$, where $\mathcal{A}^i$, a^i , and e^i denote the sets of predicted affiliations, authors, and emails.

In addition to structural mismatches, we define a parsing failure indicator δ_{parse}^i , which is set to 1 if the LLM output for document d cannot be parsed into the expected structured format (e.g., malformed delimiters, inconsistent field lengths, or invalid JSON), and 0 otherwise.

The overall success rate across the corpus $\mathcal{D}$ is then defined as:

$$\text{Success Rate} = 1 - \frac{1}{|\mathcal{D}|} \sum_{i=0}^{|\mathcal{D}|} (\delta_{\text{aff}}^i \vee \delta_{\text{email}}^i \vee \delta_{\text{parse}}^i)$$

This metric captures the fraction of documents for which both affiliation and email counts are structurally aligned with the author list, which is a necessary prerequisite for reliable metadata extraction.

Evaluation Given the scale of this benchmark, where token consumption can easily exceed the billion-token level (10^9), we employ a suite of efficient, batchable, and lower-cost pretrained LLMs that also achieve performance comparable to state-of-the-art proprietary models. Specifically, we evaluate four pretrained models from the GLM family: *glm-4*, *glm-4-plus*, *glm-4-air*, *glm-4-flash*, and *glm-3-turbo*. These models offer a strong trade-off between accuracy and throughput, making them well-suited for large-scale structured extraction tasks.

We conduct the evaluation across seven major machine learning venues over the past five years, covering more than 70,000 papers, with each experiment processing approximately 150 million tokens. Results are summarized in Table 1, highlighting structural consistency in extracting author–affiliation–email triplets. Token usage and associated inference costs are detailed in Table 2, Token Consumption Section. We omit benchmarking *glm-4* due to its prohibitive cost—each trial with 100 million tokens exceeds \$1,000.